

# Lecture 3 R Basics: "apply" functions

Zhe Zhao

Stevens Institute of Technology

*zzhao6@stevens.edu*

September 14, 2015

# Agenda

## 1 Data Types (cont'd)

- Data Frame
- Removing NA

## 2 "Apply" Functions

- Vectorized Operation
- lapply
- sapply

## 3 Generating Random Numbers

# Import CSV

We use *read.table* or *read.csv* to read tabular data. These two functions are almost identical except their default separators.

## Example

```
> read.csv("goog.csv")
> GOOG <- read.csv("goog.csv")
# default value of header is T
> GOOG <- read.csv(file = "goog.csv", header = T)

> head(GOOG) # first several rows of GOOG

# GOOG is a list
> mode(GOOG)
> names(GOOG)
> GOOG$Open
> GOOG$Adj.Close
```

# Data Frame

Data frames are used to store tabular data.

- A special type of list.
- Each element of the list (which is a vector) can be thought of as a column.
- Unlike matrices, data frames can store different classes of objects in each column (just like lists); matrices must have every element be the same class.
- Data frames are usually created by calling `read.table()` or `read.csv()`.
- Can be converted to a matrix by calling `data.matrix()`

# Data Frame

We can create data frames using the build-in function **data.frame()**.

## Example (Creating data frame)

```
> # two vectors
> kids <- c("Joe", "Jill")
> ages <- c(11, 12)
> typeof(kids)
[1] "character"
> typeof(ages)
[1] "double"
>
> # create a dataframe
> df <- data.frame(kids, ages)
> df
  kids ages
1  Joe   11
2 Jill   12
```

# Data Frame

Data frames are essentially lists.

## Example (names in data frame)

```
> # create a dataframe
> df <- data.frame(kids, ages)
> # create a list
> df2 <- list(kids, ages)
>
> # both are list
> typeof(df)
[1] "list"
> typeof(df2)
[1] "list"
```

# Data Frame

Argument *stringsAsFactors*.

Factors in R represent categorical variables, that is for statistical analysis. In finance, this data structure are not frequently used.

## Example

```
> # strings as factors
> df <- data.frame(kids, ages)
> df$kids                                # factors
[1] Joe  Jill
Levels: Jill Joe
>
> df <- data.frame(kids, ages, stringsAsFactors = F)
> df$kids                                # strings
[1] "Joe"  "Jill"
```

# Data Frame

Create a data frame by reading a csv file.

## Example

```
> aapl <- read.csv("AAPL.csv", stringsAsFactors = F)
> head(aapl)
```

	Date	Open	High	Low	Close	Volume	Adj.Close
1	2015-09-11	111.79	114.21	111.76	114.21	49441800	114.21
2	2015-09-10	110.27	113.28	109.90	112.57	62675200	112.57
3	2015-09-09	113.76	114.02	109.77	110.15	84344400	110.15
4	2015-09-08	111.75	112.56	110.32	112.31	54114200	112.31
5	2015-09-04	108.97	110.45	108.51	109.27	49963900	109.27
6	2015-09-03	112.49	112.78	110.04	110.37	52906400	110.37

```
> typeof(aapl)
[1] "list"
```



# Missing Values

Missing values are denoted by **NA** or **NaN** for undefined mathematical operations.

- *is.na()* is used to test objects if they are NA.
- *is.nan()* is used to test for NaN.
- **NA** values have a class also, so there are integer NA, character NA, etc.
- A **NaN** value is also **NA** but the converse is not true.

# Removing NA

A common task is to remove missing values from your data.

## Example

```
> x <- c(1, 2, 3, NA, NA, 6, NA, 8)
> xna <- is.na(x)      # vectorized operation
> xna
[1] FALSE FALSE FALSE  TRUE  TRUE FALSE  TRUE FALSE
> x[!xna]
[1] 1 2 3 6 8
```

# Vectorized Operation

Suppose we have a function  $f()$  that we wish to apply to all elements of a vector  $x$ . Instead of looping all elements in  $x$  and calling  $f()$  in each iteration, we can simply call  $f()$  on  $x$  itself.

This can really **simplify our code** and, moreover, give us **a dramatic performance increase** of hundredsfold or more.

# Vectorized Operation

## Example

```
> x <- 1:4
> y <- 6:9
> # want to do x + y

> # using for loop
> result <- c()
> for (i in x)
+ {
+   result[i] = x[i] + y[i]
+ }

> result
[1] 7 9 11 13
```

# Vectorized Operation

We can directly apply "+" to x and y.

## Example

```
> result <- x + y
> result
[1] 7 9 11 13
> rbind(x, y, result)
      [,1] [,2] [,3] [,4]
x         1    2    3    4
y         6    7    8    9
result     7    9   11   13
```

# Vectorized Operation

Another example

## Example

```
> x <- 1:4  
> x > 2  
[1] FALSE FALSE  TRUE  TRUE
```

# Vectorized Operation

Apply vectorized operation to functions defined by user

## Example

```
> # user defined functions  
> mySquare <- function(x)  
+ {  
+   return (x^2)  
+ }
```

# Vectorized Operation

## Example

```
> x <- 1:4
> y <- 6:9
> mySquare(x)
[1] 1 4 9 16
> mySquare(y)
[1] 36 49 64 81
> table1 <- cbind(x, mySquare(x), y, mySquare(y))
> table1
```

	x		y	
[1,]	1	1	6	36
[2,]	2	4	7	49
[3,]	3	9	8	64
[4,]	4	16	9	81



# Vectorized Operation

Revisiting the coin flipping example.

## Example

```
No.heads <- 0
result.Vec <- NULL
for (flips in 1:1000)
{
  # x = c(1, 0), 1 means head, 0 means tail
  tmp <- sample(x=c(1, 0), size=1,
               replace=T, prob=c(0.5, 0.5))
  No.heads <- No.heads + tmp
  result.Vec <- c(result.Vec, No.heads/flips)
}
plot(1:1000, result.Vec, type="l")      # produce a figure
```

# Vectorized Operation

Rewrite the same example in vectors

## Example

```
> ?cumsum # please check by yourself
>
> sims <- sample(c(1, 0), size = 1000,
                 replace = T, prob = c(0.5, 0.5))
> cum_sims <- cumsum(sims)           # cumulative Sums
> result.Vec <- cum_sims / (1:1000)  # vectorized operation
>
> plot.ts(result.Vec)
```

# “Apply” Functions

## **lapply**

- *lapply* takes two arguments: a list *x*, a function or a name of function.
- If *x* is not a list, it will be coerced to a list using `as.list()`.
- *lapply* always returns a list, regardless of the class of the input.

## Example

```
> x <- 1:4  
> lapply(x, mySquare)  
[[1]]  
[1] 1  
  
[[2]]  
[1] 4  
  
[[3]]  
[1] 9  
  
[[4]]  
[1] 16
```

*sapply* will try to simplify the result of *lapply* if possible.

- If the result is a list where every element is length 1, then a vector is returned.
- If the result is a list where every element is a vector of the same length, a matrix is returned.
- Otherwise a list is returned.

## Example

```
> x <- 1:4  
> sapply(x, mySquare)  
[1] 1 4 9 16
```

# Apply Functions

## Another example

### Example

```
> x <- list(rnorm(10000), runif(10000, min = 0, max = 1))
> lapply(x, mean)
[[1]]
[1] -0.008742473

[[2]]
[1] 0.5009495

> sapply(x, mean)
[1] -0.008742473 0.500949453
```

# Apply Functions

## More example

### Example

```
> x <- 1:4  
> lapply(x, runif)  
[[1]]  
[1] 0.8414947  
  
[[2]]  
[1] 0.6263586 0.8831707  
  
[[3]]  
[1] 0.6218396 0.1194392 0.3133257  
  
[[4]]  
[1] 0.8740655 0.4518131 0.1776370 0.5959832
```



# About Performance

Use "apply" functions will dramatically increase the performance of your code

## Example

```
> x <- 1:10^5
> system.time(
+   for (i in x)
+   {
+       mySquare(x)
+   }
+ )
   user  system elapsed 
44.170   1.718  45.880
```

# About Performance

Here comes the magic.

## Example

```
> system.time(  
+   sapply(x, mySquare)  
+ )  
   user  system elapsed  
0.175   0.002   0.176
```

# Other Apply Functions

- `lapply`      Apply a Function over a List or Vector
- `sapply`      Simplify the result of `lapply`
  
- `apply`      Apply Functions Over Array Margins
- `eapply`      Apply a Function Over Values in an Environment
- `mapply`      Apply a Function to Multiple List or Vector Arguments
- `rapply`      Recursively Apply a Function to a List
- `tapply`      Apply a Function Over a Ragged Array

# Normal Distribution


Now let's move on how to generate random variables.

## Example

```
> # rnorm(n = , mean = , sd = )  
  
> x <- rnorm(n = 10000, mean = 0, sd = 1)  
> hist(x)  
> hist(x, nclass = 40)  
  
> # another sigma  
> x <- rnorm(n = 10000, mean = 0, sd = 5)
```


# Normal Distribution

**$\mu = 0$ ,  $\sigma = 1$**



# Normal Distribution

**$\mu = 0$ ,  $\sigma = 5$**



# Normal Distribution

## **set.seed()**

Anywho the random numbers R gives you aren't really random. They're pseudo-random. Basically there's a function that outputs numbers that look random. To do this it needs some inputs. The first input it gets will be the 'seed'.

### Example

```
> set.seed(1)
> rnorm(5)
[1] -0.6264538  0.1836433 -0.8356286  1.5952808  0.3295078
> rnorm(5)
[1] -0.8204684  0.4874291  0.7383247  0.5757814 -0.3053884
> set.seed(1)
> rnorm(5)
[1] -0.6264538  0.1836433 -0.8356286  1.5952808  0.3295078
```

# Normal Distribution

- `rnorm`: generate random Normal variates with a given mean and standard deviation.
- `dnorm`: evaluate the Normal probability density (with a given mean/SD) at a point (or vector of points)
- `pnorm`: evaluate the cumulative distribution function for a Normal distribution
- `qnorm`: gives the quantile function



# Normal Distribution

## Example

```
> # Density function
> dnorm(x = 0)      # mean = 0, sd = 1
[1] 0.3989423
> dnorm(x = 1)      # check table if you want
[1] 0.2419707

> # cumulative distribution function
> pnorm(q = 0)
[1] 0.5
> pnorm(q = 5)
[1] 0.9999997
```

# Other Distributions

- `rt()` – t distribution
- `rpois()` – poisson distribution
- `runif()` – uniform distribution
- `rexp()` – exponential distribution

## Example

```
> x <- rpois(1000, lambda = 2)
> hist(x, nclass = 40)
>
> x <- rexp(1000)
> hist(x, nclass = 40)
>
> x <- rt(1000, df = 10)
> hist(x, nclass = 40)
```